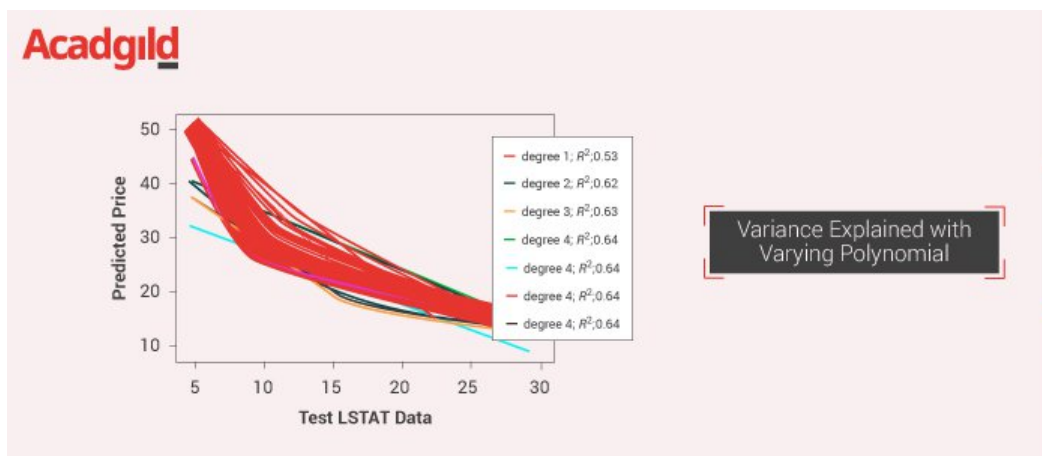


Understand Power of Polynomials with Polynomial Regression



Abhay Kumar • September 13, 2018 1 3,527




This entry is part 4 of 17 in the series [Machine Learning Algorithms](#)

Introduction to Polynomial Regression

Polynomial regression is a special case of linear regression. It's based on the idea of how to your select your features. Looking at the multivariate regression with 2 variables: x_1 and x_2 . Linear regression will look like this:

$$y = a_1 * x_1 + a_2 * x_2.$$



Free Step-by-step Guide To Become A Data Scientist

Subscribe and get this detailed guide absolutely FREE

Download Now!

Now you want to have a polynomial regression (let's make 2-degree polynomial). We will create a few additional features: $x_1 * x_2$, x_1^2 and x_2^2 . So we will get your 'linear regression':

$$y = a_1 * x_1 + a_2 * x_2 + a_3 * x_1 * x_2 + a_4 * x_1^2 + a_5 * x_2^2$$

A polynomial term: a quadratic (squared) or cubic (cubed) term turns a linear regression model into a curve. But because it is the data X that is squared or cubed, not the Beta coefficient, it still qualifies as a linear model.

This makes it a nice and straightforward way to model curves without having to model complicated nonlinear models.

One common pattern within machine learning is to use linear models trained on nonlinear functions of the data. This approach maintains the generally fast performance of linear methods while allowing them to fit a much wider range of data.

For example, a simple linear regression can be extended by constructing polynomial features from the coefficients. In the standard linear regression case, you might have a model that looks like this for two-dimensional data:

$$\hat{y}(w, x) = w_0 + w_1 x_1 + w_2 x_2$$

If we want to fit a paraboloid to the data instead of a plane, we can combine the features in second-order polynomials, so that the model looks like this:

$$\hat{y}(w, x) = w_0 + w_1x_1 + w_2x_2 + w_3x_1x_2 + w_4x_1^2 + w_5x_2^2$$

The (sometimes surprising) observation is that this is still a linear model: to see this, imagine creating a new variable

$$z = [x_1, x_2, x_1x_2, x_1^2, x_2^2]$$

With this re-labeling of the data, our problem can be written

$$\hat{y}(w, x) = w_0 + w_1z_1 + w_2z_2 + w_3z_3 + w_4z_4 + w_5z_5$$

We see that the resulting polynomial regression is in the same class of linear models we'd considered above (i.e. the model is linear in w) and can be solved by the same techniques.

By considering, linear fits within a higher-dimensional space built with these basis functions, the model has the flexibility to fit a much broader range of data.

Here is an example of applying this idea to one-dimensional data, using polynomial features of varying degrees:

```
from sklearn.preprocessing import PolynomialFeatures
import numpy as np
X = np.arange(6).reshape(3, 2)
X

poly = PolynomialFeatures(degree=2)
poly.fit_transform(X)
```

In some cases, it's not necessary to include higher powers of any single feature, but only the so-called interaction features that multiply together at most d distinct features. These can be derived from `PolynomialFeatures` with the setting `interaction_only=True`.

For example, when dealing with boolean features, $x_i^n = x_i$ for all n , and is therefore useless; but $x_i x_j$ represents the conjunction of two booleans. This way, we can solve the XOR problem with a linear classifier:

```
from sklearn.linear_model import Perceptron
from sklearn.preprocessing import PolynomialFeatures
import numpy as np
X = np.array([[0, 0], [0, 1], [1, 0], [1, 1]])
y = X[:, 0] ^ X[:, 1]
y

X = PolynomialFeatures(interaction_only=True).fit_transform(X).astype(float)
X
```

Problem Statement:

To understand the relationship between the degree of the independent variable and the dependent variable.

Data details

```
Boston House Prices dataset
=====

Notes
-----

Data Set Characteristics:

    :Number of Instances: 506

    :Number of Attributes: 13 numeric/categorical predictive

    :Median Value (attribute 14) is usually the target

    :Attribute Information (in order):
        - CRIM    per capita crime rate by town
        - ZN      proportion of residential land zoned for lots over 25,000 sq. ft.
        - INDUS   proportion of non-retail business acres per town
        - CHAS    Charles River dummy variable (= 1 if tract bounds river; 0 otherwise)
        - NOX     nitric oxides concentration (parts per 10 million)
        - RM      average number of rooms per dwelling
        - AGE     proportion of owner-occupied units built prior to 1940
        - DIS     weighted distances to five Boston employment centres
```

- RAD index of accessibility to radial highways
- TAX full-value property-tax rate per \$10,000
- PTRATIO pupil-teacher ratio by town
- B $1000(B_k - 0.63)^2$ where B_k is the proportion of blacks living in the town
- LSTAT - MEDV Median value of owner-occupied homes in the town

:Missing Attribute Values: None
 % lower status of the population

:Creator: Harrison, D. and Rubinfeld, D.L.

This is a copy of UCI ML housing dataset.

<http://archive.ics.uci.edu/ml/datasets/Housing>

This dataset was taken from the StatLib library which is maintained at

The Boston house-price data of Harrison, D. and Rubinfeld, D.L. 'Hedonic prices and the demand for clean air', J. Environ. Economics & Management vol.5, 81-102, 1978. Used in Belsley, Kuh & Welsch, 'Regression diagnostics...', Wiley, 1980. N.B. Various transformations are used in the tables pages 244-261 of the latter.

The Boston house-price data has been used in many machine learning papers that address regression problems.

Tools used :

- Pandas
- Numpy
- Matplotlib
- scikit-learn

Import necessary libraries

Import the necessary modules from specific libraries.

```
import numpy as np
import pandas as pd
%matplotlib inline
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn import datasets
from sklearn.metrics import mean_squared_error
```

```
from sklearn.preprocessing import PolynomialFeatures
from sklearn.linear_model import LinearRegression
from sklearn.pipeline import make_pipeline
```

Load the data set

Use the pandas module to read the taxi data from the file system. Check few records of the dataset.

```
#
#####
# Load data
boston = datasets.load_boston()
print(boston.data.shape, boston.target.shape)
print(boston.feature_names)

(506, 13) (506,)
['CRIM' 'ZN' 'INDUS' 'CHAS' 'NOX' 'RM' 'AGE' 'DIS' 'RAD' 'TAX' 'PTRATIO' 'B
```

```
data = pd.DataFrame(boston.data, columns=boston.feature_names)
data = pd.concat([data, pd.Series(boston.target, name='MEDV')], axis=1)
data.head()

   CRIM    ZN  INDUS  CHAS  NOX   RM   AGE  DIS   RAD  TAX  PTRATIO  B
0  0.00632  18.0   2.31   0.0  0.538  6.575  65.2  4.0900  1.0  296.0   15.3   35
1  0.02731   0.0   7.07   0.0  0.469  6.421  78.9  4.9671  2.0  242.0   17.8   39
2  0.02729   0.0   7.07   0.0  0.469  7.185  61.1  4.9671  2.0  242.0   17.8   39
3  0.03237   0.0   2.18   0.0  0.458  6.998  45.8  6.0622  3.0  222.0   18.7   39
4  0.06905   0.0   2.18   0.0  0.458  7.147  54.2  6.0622  3.0  222.0   18.7   39
```

Feature Selection:

Let's select one of the numerical fields for model fitting here "LSTAT" i.e. % of the population in the neighborhood which comes under lower economic status.

```
X = data[['LSTAT']]
y = data['MEDV']
```

Train test split :

```
x_training_set, x_test_set, y_training_set, y_test_set = train_test_split
```

Training/model fitting with various degrees :

Fit the model to selected supervised data. We will use the API called Polynomial Features which takes the parameter as the degree of the polynomial. With the given polynomial degree we will fit the data with the linear regression model.

The motive of this fitting is to see if there is a better explanation of the variance with an increase in the degree of the polynomial of the selected feature.

```
# Polynomial Regression-nth order
plt.scatter(x_test_set, y_test_set, s=10, alpha=0.3)

for degree in [1,2,3,4,5,6,7]:

    model = make_pipeline(PolynomialFeatures(degree), LinearRegression())

    model.fit(x_training_set,y_training_set)

    y_plot = model.predict(x_test_set)

    plt.plot(x_test_set, y_plot, label="degree %d" % degree

            +'; $R^2$: %.2f' % model.score(x_test_set, y_test_set))

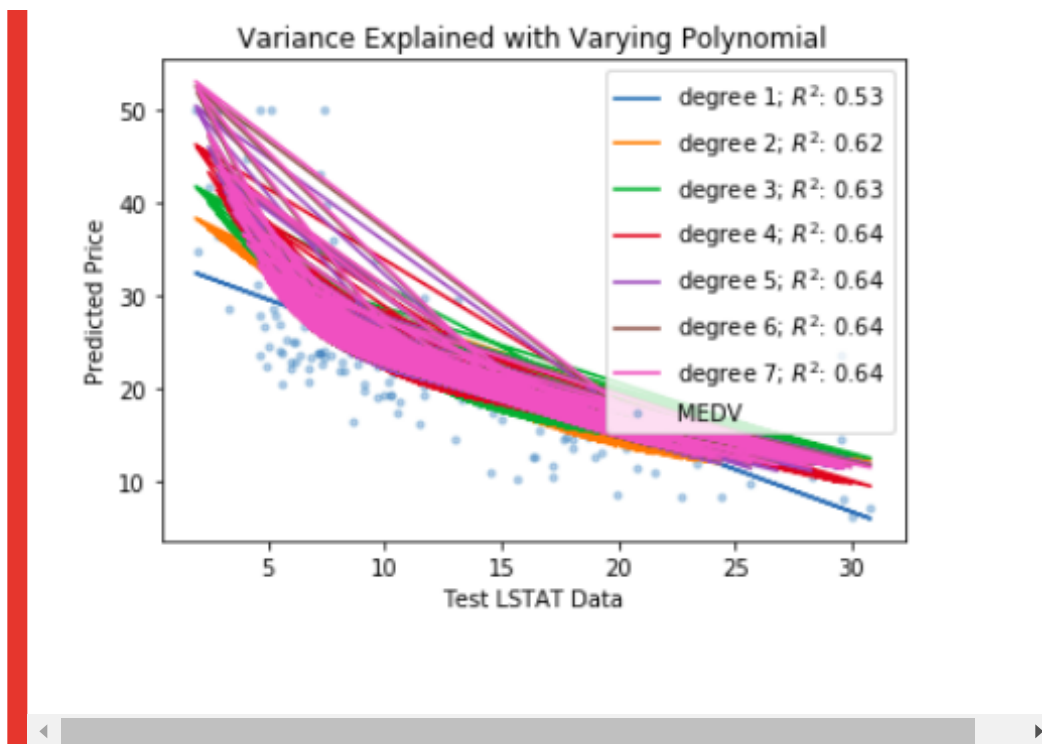
plt.legend(loc='upper right')

plt.xlabel("Test LSTAT Data")

plt.ylabel("Predicted Price")

plt.title("Variance Explained with Varying Polynomial")

plt.show()
```




We can notice that R^2 has increased from 0.53 \rightarrow 0.62 \rightarrow 0.63 \rightarrow 0.64 with rising in the degree of the polynomial initially and then it has saturated at 0.64.

Series Navigation

< < Covariance and Correlation

Logistic Regression With PCA – Speeding Up and Benchmarking > >


Related



Logistic Regression

July 12, 2018


In "Data Science and Artificial Intelligence"



Multiple Linear Regression

September 14, 2018

In "Data Science and Artificial Intelligence"



Multivariate Multilabel Classification with Logistic Regression

September 12, 2018

In "Data Science and Artificial Intelligence"

One Comment

Pingback: Polynomial Regression in Machine Learning with Example - IntelliNova

This site uses Akismet to reduce spam. [Learn how your comment data is processed.](#)

