

# Nearest Neighbor Classification

Machine Learning  
Fall 2017



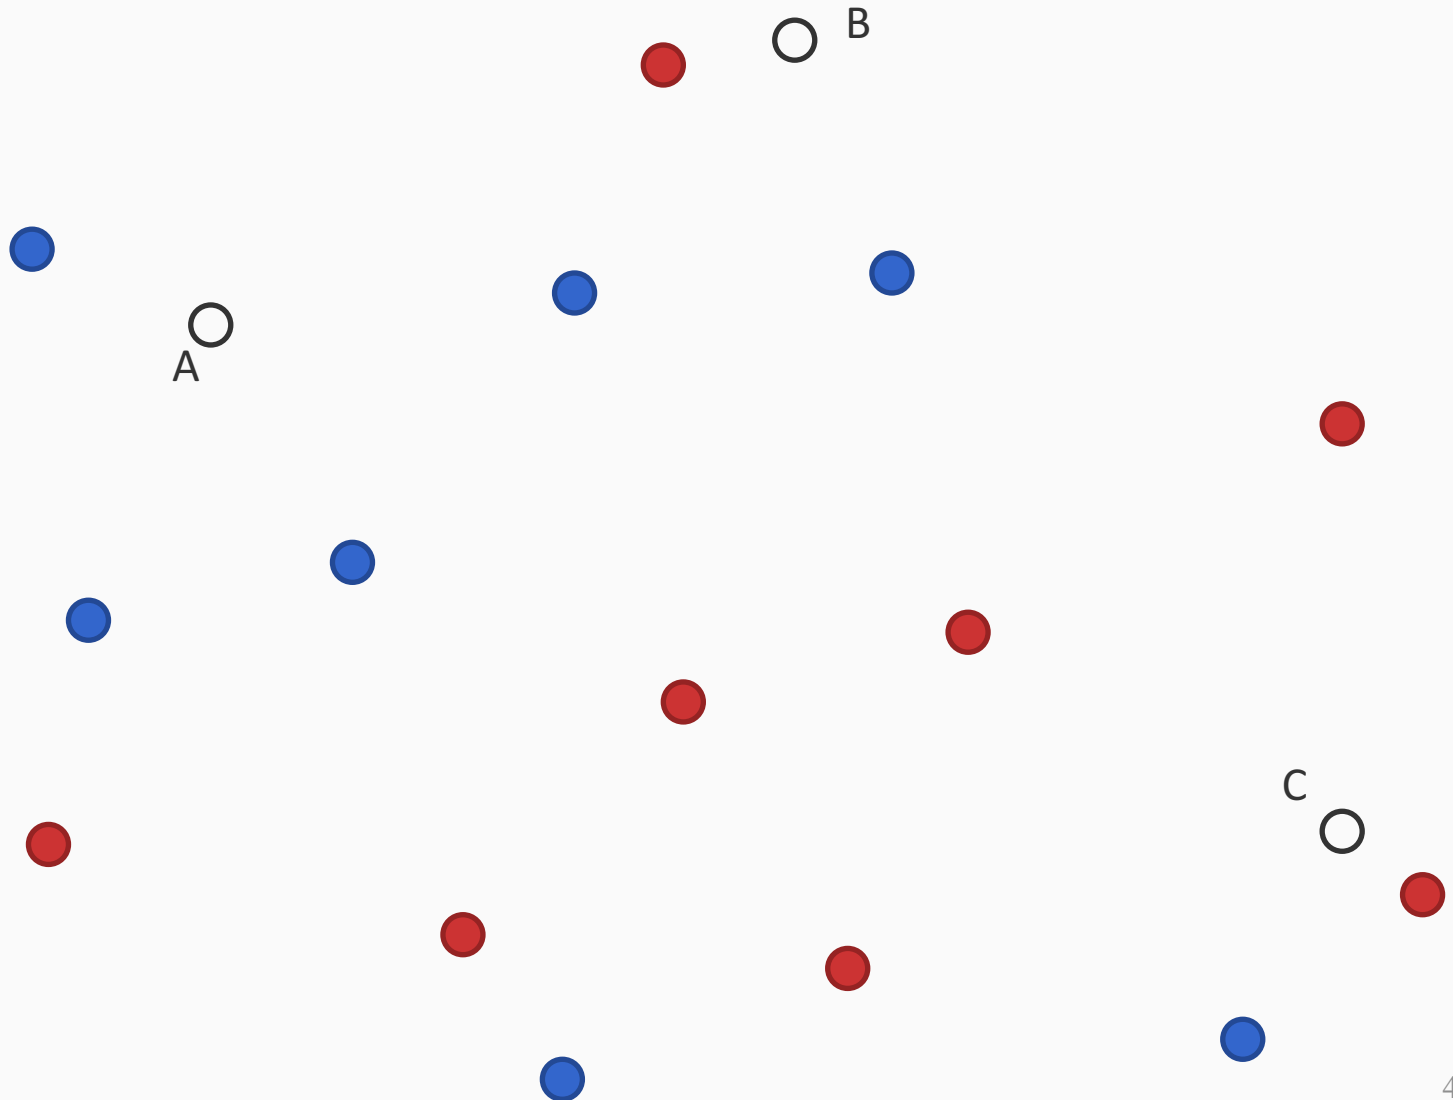
# This lecture

- K-nearest neighbor classification
  - The basic algorithm
  - Different distance measures
  - Some practical aspects
- Voronoi Diagrams and Decision Boundaries
  - What is the hypothesis space?
- The Curse of Dimensionality

# This lecture

- K-nearest neighbor classification
  - The basic algorithm
  - Different distance measures
  - Some practical aspects
- Voronoi Diagrams and Decision Boundaries
  - What is the hypothesis space?
- The Curse of Dimensionality

# How would you color the blank circles?



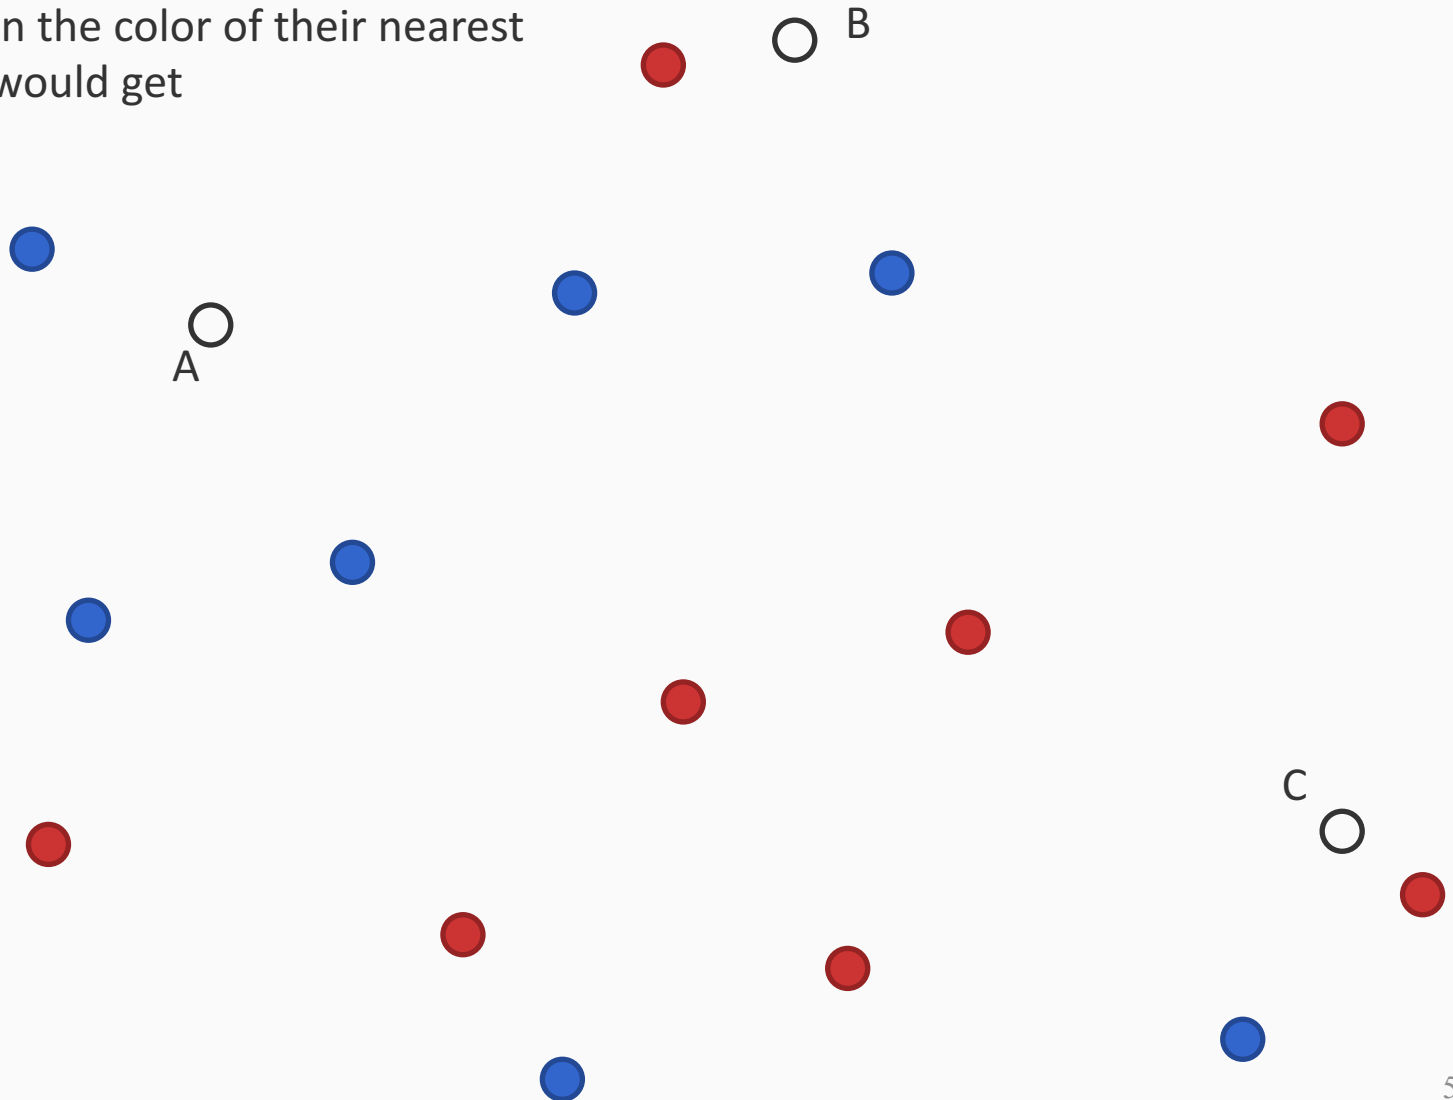
# How would you color the blank circles?

If we based it on the color of their nearest neighbors, we would get

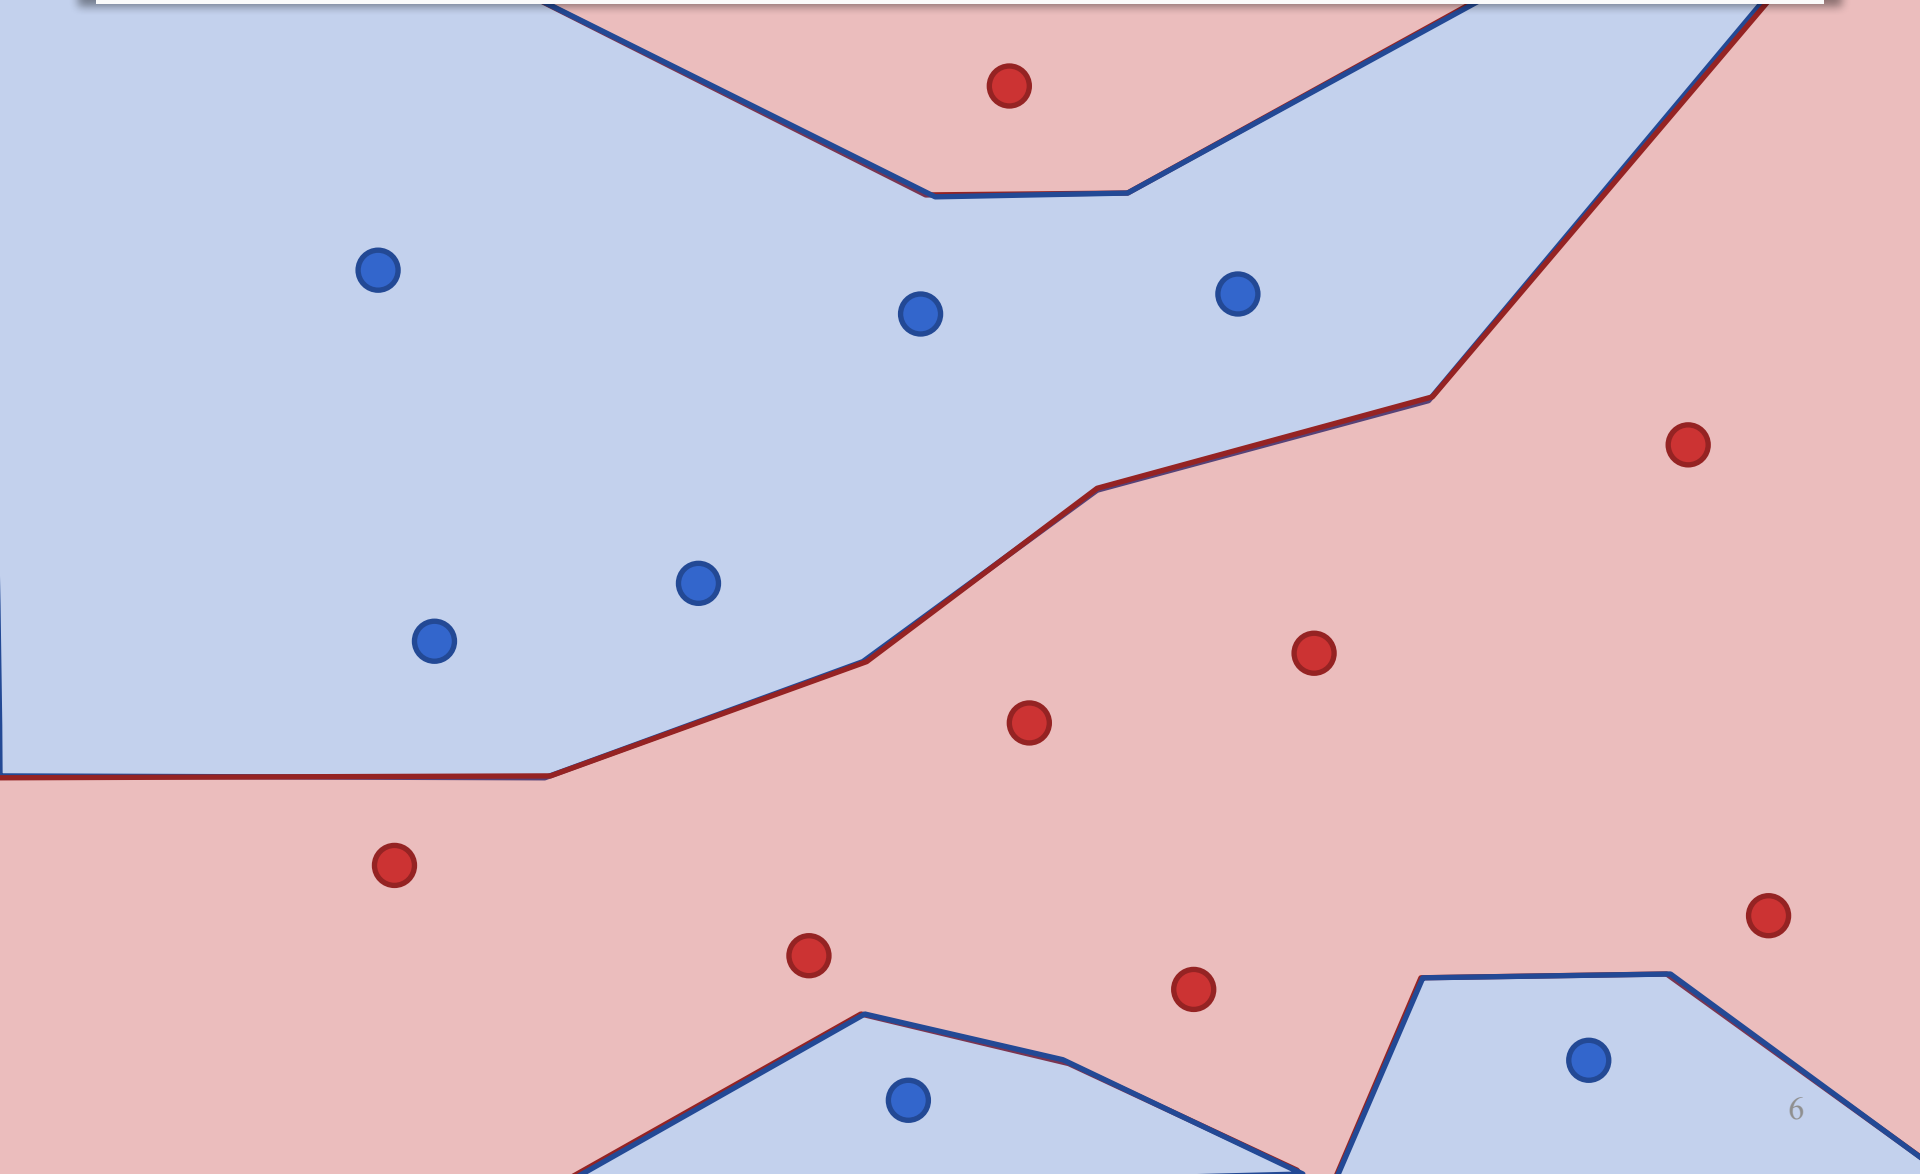
A: Blue

B: Red

C: Red



Training data partitions the entire instance space  
(using labels of nearest neighbors)



# Nearest Neighbors: The basic version

- Training examples are vectors  $\mathbf{x}_i$  associated with a label  $y_i$ 
  - E.g.  $\mathbf{x}_i$  = a feature vector for an email,  $y_i$  = SPAM
- Learning: Just store all the training examples
- Prediction for a new example  $\mathbf{x}$ 
  - Find the training example  $\mathbf{x}_i$  that is *closest* to  $\mathbf{x}$
  - Predict the label of  $\mathbf{x}$  to the label  $y_i$  associated with  $\mathbf{x}_i$

# K-Nearest Neighbors

- Training examples are vectors  $\mathbf{x}_i$  associated with a label  $y_i$ 
  - E.g.  $\mathbf{x}_i$  = a feature vector for an email,  $y_i$  = SPAM
- Learning: Just store all the training examples
- Prediction for a new example  $\mathbf{x}$ 
  - Find the  $k$  *closest* training examples to  $\mathbf{x}$
  - Construct the label of  $\mathbf{x}$  using these  $k$  points. *How?*
  - For classification: ?



# K-Nearest Neighbors

- Training examples are vectors  $\mathbf{x}_i$  associated with a label  $y_i$ 
  - E.g.  $\mathbf{x}_i$  = a feature vector for an email,  $y_i$  = SPAM
- Learning: Just store all the training examples
- Prediction for a new example  $\mathbf{x}$ 
  - Find the  $k$  *closest* training examples to  $\mathbf{x}$
  - Construct the label of  $\mathbf{x}$  using these  $k$  points. *How?*
  - For classification: Every neighbor votes on the label. Predict the most frequent label among the neighbors.
  - For regression: ?

# K-Nearest Neighbors

- Training examples are vectors  $\mathbf{x}_i$  associated with a label  $y_i$ 
  - E.g.  $\mathbf{x}_i$  = a feature vector for an email,  $y_i$  = SPAM
- Learning: Just store all the training examples
- Prediction for a new example  $\mathbf{x}$ 
  - Find the  $k$  *closest* training examples to  $\mathbf{x}$
  - Construct the label of  $\mathbf{x}$  using these  $k$  points. *How?*
  - For classification: Every neighbor votes on the label. Predict the most frequent label among the neighbors.
  - For regression: Predict the mean value

# Instance based learning

- A class of learning methods
  - Learning: Storing examples with labels
  - Prediction: When presented a new example, classify the labels using *similar* stored examples
- K-nearest neighbors algorithm is an example of this class of methods
- Also called *lazy* learning, because most of the computation (in the simplest case, all computation) is performed only at prediction time

# Distance between instances

- In general, a good place to inject knowledge about the domain
- Behavior of this approach can depend on this
- How do we measure distances between instances?

# Distance between instances

Numeric features, represented as  $n$  dimensional vectors

# Distance between instances

Numeric features, represented as n dimensional vectors

- Euclidean distance

$$\|\mathbf{x}_1 - \mathbf{x}_2\|_2 = \sqrt{\sum_{i=1}^n (\mathbf{x}_{1,i} - \mathbf{x}_{2,i})^2}$$

# Distance between instances

Numeric features, represented as  $n$  dimensional vectors

- Euclidean distance

$$\|\mathbf{x}_1 - \mathbf{x}_2\|_2 = \sqrt{\sum_{i=1}^n (\mathbf{x}_{1,i} - \mathbf{x}_{2,i})^2}$$

- Manhattan distance

$$\|\mathbf{x}_1 - \mathbf{x}_2\|_1 = \sum_{i=1}^n |\mathbf{x}_{1,i} - \mathbf{x}_{2,i}|$$

# Distance between instances

Numeric features, represented as  $n$  dimensional vectors

- Euclidean distance

$$\|\mathbf{x}_1 - \mathbf{x}_2\|_2 = \sqrt{\sum_{i=1}^n (\mathbf{x}_{1,i} - \mathbf{x}_{2,i})^2}$$

- Manhattan distance

$$\|\mathbf{x}_1 - \mathbf{x}_2\|_1 = \sum_{i=1}^n |\mathbf{x}_{1,i} - \mathbf{x}_{2,i}|$$

- $L_p$ -norm

- Euclidean =  $L_2$
- Manhattan =  $L_1$
- Exercise: What is  $L_\infty$ ?

$$\|\mathbf{x}_1 - \mathbf{x}_2\|_p = \left( \sum_{i=1}^n |\mathbf{x}_{1,i} - \mathbf{x}_{2,i}|^p \right)^{\frac{1}{p}}$$



# Distance between instances

What about symbolic/categorical features?

# Distance between instances

## Symbolic/categorical features

Most common distance is the *Hamming distance*

- Number of bits that are different
- Or: Number of features that have a different value
- Also called the *overlap*
- Example:

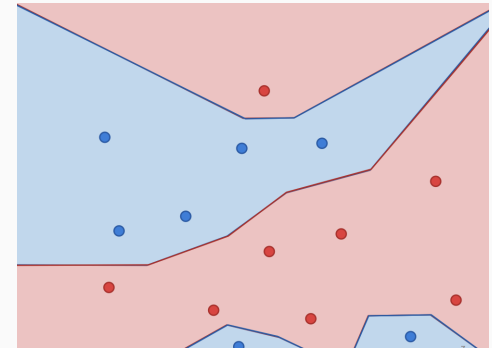
$X_1$ : {Shape=Triangle, Color=Red, Location=Left, Orientation=Up}

$X_2$ : {Shape=Triangle, Color=Blue, Location=Left, Orientation=Down}

Hamming distance = 2

# Advantages

- Training is *very fast*
  - Just adding labeled instances to a list
  - More complex indexing methods can be used, which slow down learning slightly to make prediction faster
- Can learn very *complex functions*
- We always have the training data
  - For other learning algorithms, after training, we don't store the data anymore. What if we want to do something with it later...



# Disadvantages

- Needs a lot of storage
  - Is this really a problem now?
- Prediction can be slow!
  - Naïvely:  $O(dN)$  for  $N$  training examples in  $d$  dimensions
  - More data will make it slower
  - Compare to other classifiers, where prediction is very fast
- Nearest neighbors are fooled by irrelevant attributes
  - Important and subtle

Questions?

# Summary: K-Nearest Neighbors

- Probably the first “machine learning” algorithm
  - Guarantee: If there are enough training examples, the error of the nearest neighbor classifier will converge to the error of the optimal (i.e. best possible) predictor
- In practice, use an odd K. Why?

# Summary: K-Nearest Neighbors

- Probably the first “machine learning” algorithm
  - Guarantee: If there are enough training examples, the error of the nearest neighbor classifier will converge to the error of the optimal (i.e. best possible) predictor
- In practice, use an odd K. Why?
  - To break ties

# Summary: K-Nearest Neighbors

- Probably the first “machine learning” algorithm
  - Guarantee: If there are enough training examples, the error of the nearest neighbor classifier will converge to the error of the optimal (i.e. best possible) predictor
- In practice, use an odd K. Why?
  - To break ties
- How to choose K? Using a held-out set or by cross-validation

# Summary: K-Nearest Neighbors

- Probably the first “machine learning” algorithm
  - Guarantee: If there are enough training examples, the error of the nearest neighbor classifier will converge to the error of the optimal (i.e. best possible) predictor
- In practice, use an odd K. Why?
  - To break ties
- How to choose K? Using a held-out set or by cross-validation
- Feature normalization could be important
  - Often, good idea to center the features to make them zero mean and unit standard deviation. Why?



# Summary: K-Nearest Neighbors

- Probably the first “machine learning” algorithm
  - Guarantee: If there are enough training examples, the error of the nearest neighbor classifier will converge to the error of the optimal (i.e. best possible) predictor
- In practice, use an odd K. Why?
  - To break ties
- How to choose K? Using a held-out set or by cross-validation
- Feature normalization could be important
  - Often, good idea to center the features to make them zero mean and unit standard deviation. Why?
  - Because different features could have different scales (weight, height, etc); but the distance weights them equally
- Variants exist
  - Neighbors’ labels could be weighted by their distance

# Where are we?

- K-nearest neighbor classification
  - The basic algorithm
  - Different distance measures
  - Some practical aspects
- Voronoi Diagrams and Decision Boundaries
  - What is the hypothesis space?
- The Curse of Dimensionality

# Where are we?

- K-nearest neighbor classification
  - The basic algorithm
  - Different distance measures
  - Some practical aspects
- Voronoi Diagrams and Decision Boundaries
  - What is the hypothesis space?
- The Curse of Dimensionality

# The decision boundary for KNN

Is the K nearest neighbors algorithm explicitly building a function?

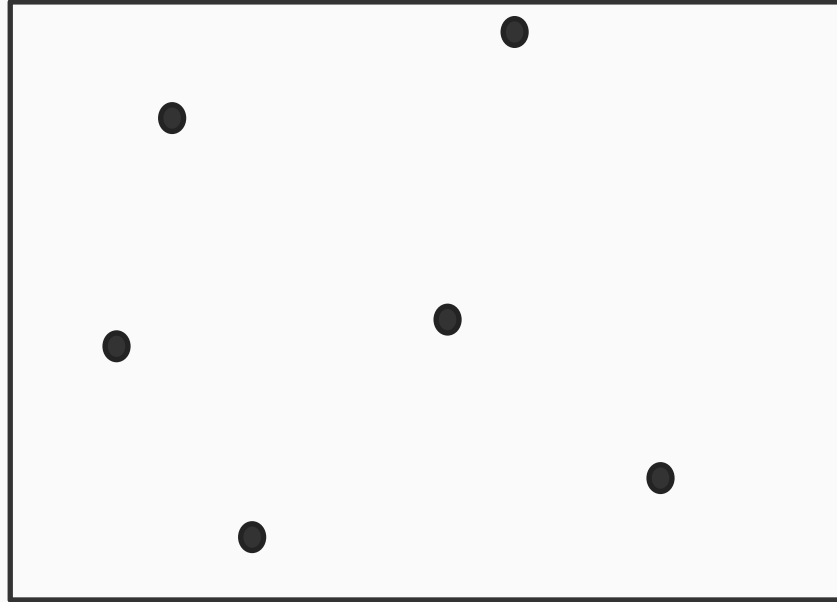
# The decision boundary for KNN

Is the K nearest neighbors algorithm explicitly building a function?

- **No**, it never forms an explicit hypothesis

But we can still ask: Given a training set what is the implicit function that is being computed

# The Voronoi Diagram

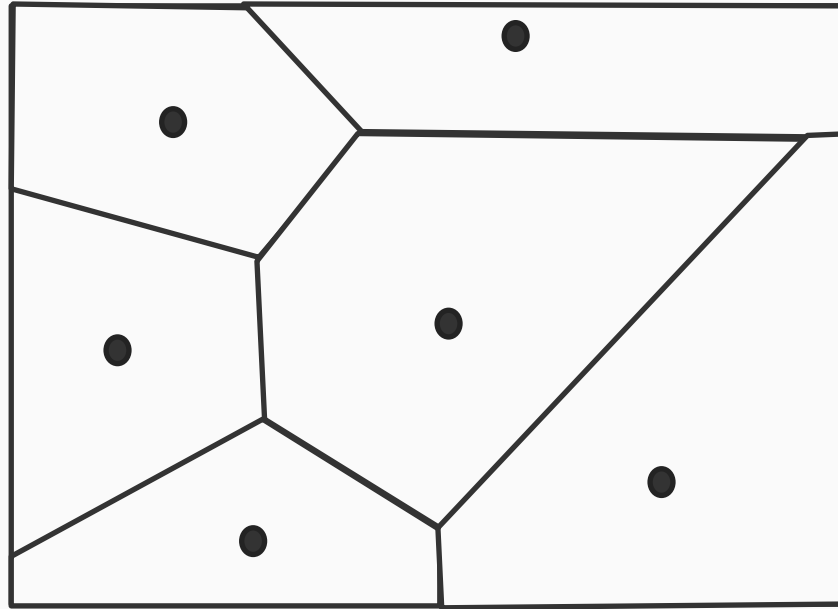


For any point  $\mathbf{x}$  in a training set  $S$ , the **Voronoi Cell** of  $\mathbf{x}$  is a polyhedron consisting of all points closer to  $\mathbf{x}$  than any other points in  $S$

The **Voronoi diagram** is the union of all Voronoi cells

- Covers the entire space

# The Voronoi Diagram

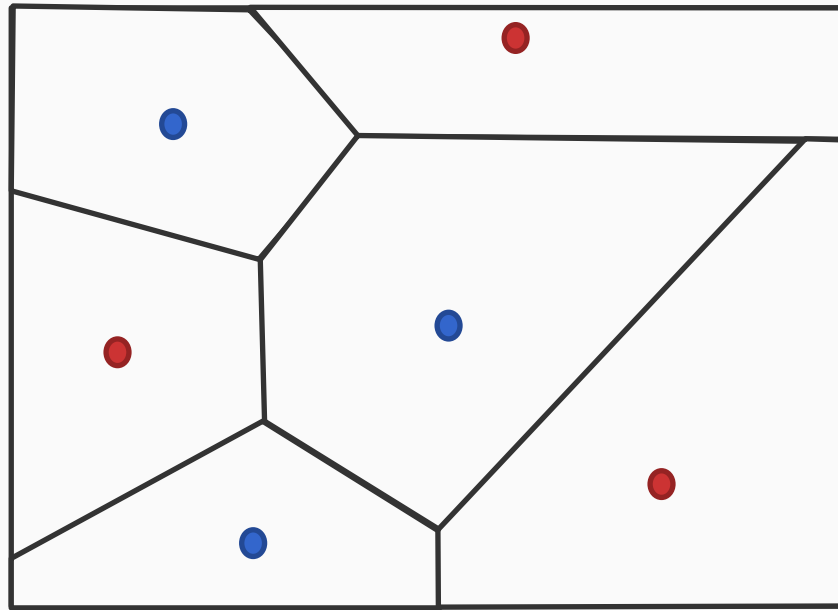


For any point  $\mathbf{x}$  in a training set  $S$ , the **Voronoi Cell** of  $\mathbf{x}$  is a polyhedron consisting of all points closer to  $\mathbf{x}$  than any other points in  $S$

The **Voronoi diagram** is the union of all Voronoi cells

- Covers the entire space

# Voronoi diagrams of training examples



Points in the Voronoi cell of a training example are closer to it than any others

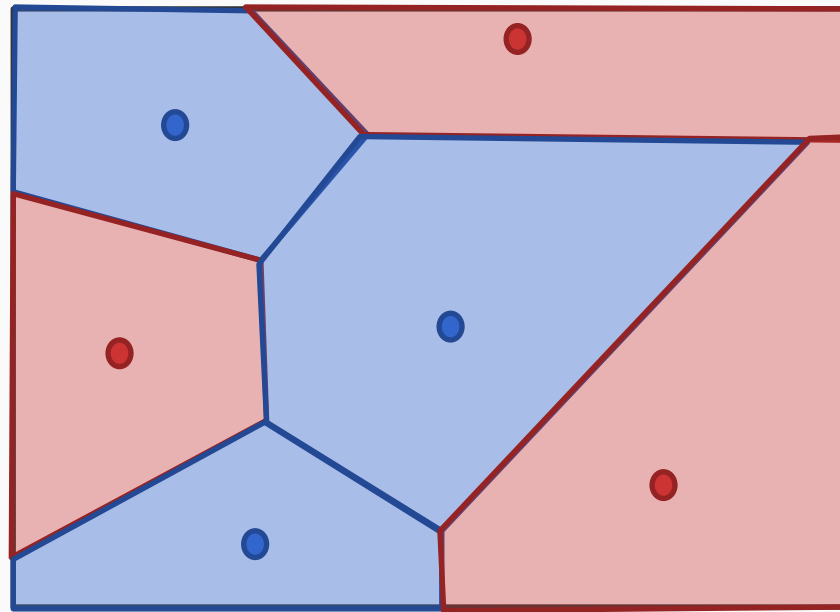
For any point  $\mathbf{x}$  in a training set  $S$ , the **Voronoi Cell** of  $\mathbf{x}$  is a polytope consisting of all points closer to  $\mathbf{x}$  than any other points in  $S$

The **Voronoi diagram** is the union of all Voronoi cells

- Covers the entire space



# Voronoi diagrams of training examples



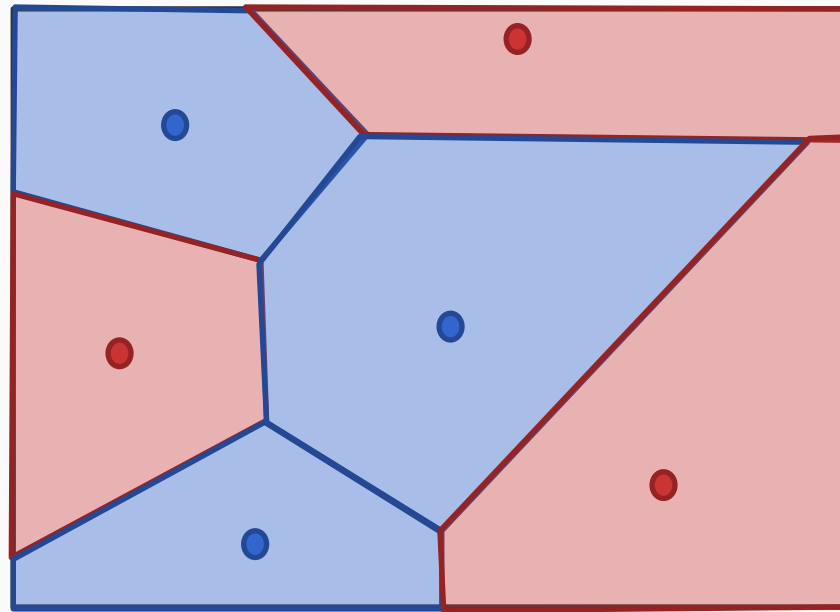
Points in the Voronoi cell of a training example are closer to it than any others

For any point  $\mathbf{x}$  in a training set  $S$ , the **Voronoi Cell** of  $\mathbf{x}$  is a polytope consisting of all points closer to  $\mathbf{x}$  than any other points in  $S$

The **Voronoi diagram** is the union of all Voronoi cells

- Covers the entire space

# Voronoi diagrams of training examples



Points in the Voronoi cell of a training example are closer to it than any others

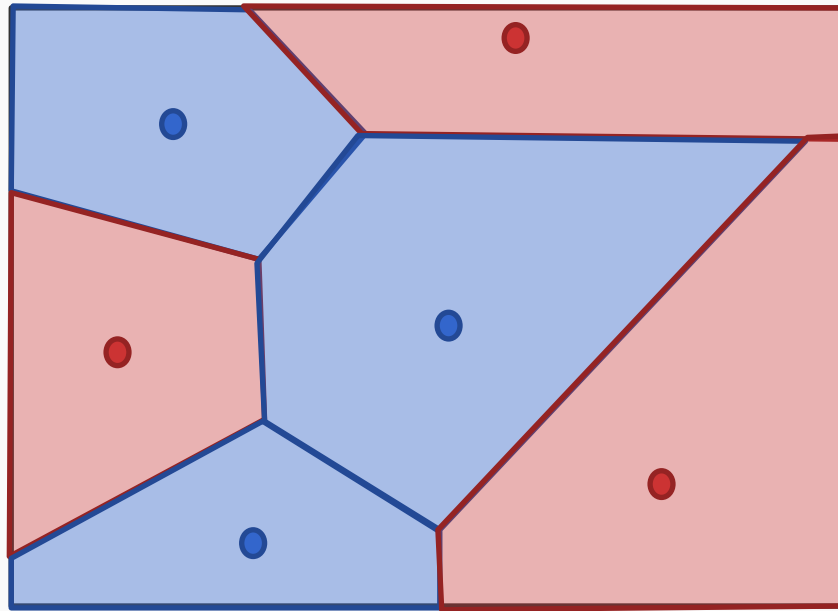
Picture uses Euclidean distance with 1-nearest neighbors.

What about K-nearest neighbors?

Also partitions the space, but much more complex decision boundary

# Voronoi diagrams of training examples

What about points on the boundary? What label will they get?



Points in the Voronoi cell of a training example are closer to it than any others

Picture uses Euclidean distance with 1-nearest neighbors.

What about K-nearest neighbors?

Also partitions the space, but much more complex decision boundary

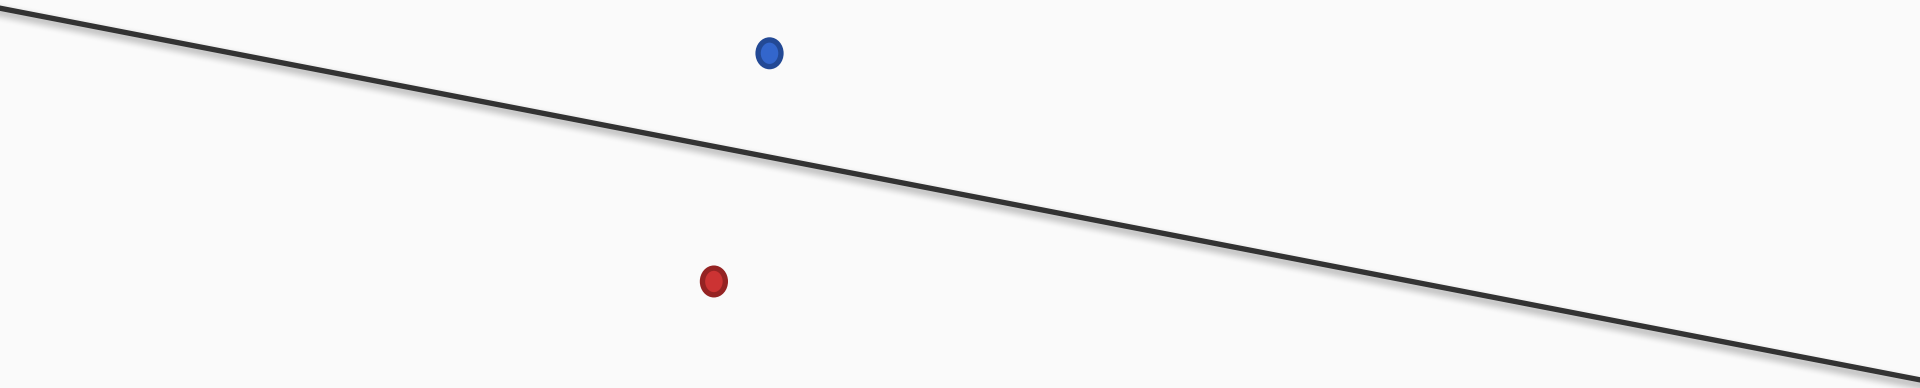
# Exercise

If you have only two training points, what will the decision boundary for 1-nearest neighbor be?



# Exercise

If you have only two training points, what will the decision boundary for 1-nearest neighbor be?



- A line bisecting the two points

# This lecture

- K-nearest neighbor classification
  - The basic algorithm
  - Different distance measures
  - Some practical aspects
- Voronoi Diagrams and Decision Boundaries
  - What is the hypothesis space?
- The Curse of Dimensionality

# Why your classifier might go wrong

Two important considerations with learning algorithms

- **Overfitting:** We have already seen this
- **The curse of dimensionality**
  - Methods that work with low dimensional spaces may fail in high dimensions
  - What might be intuitive for 2 or 3 dimensions do not always apply to high dimensional spaces

Check out the 1884 book *Flatland: A Romance of Many Dimensions* for a fun introduction to the fourth dimension

# Of course, irrelevant attributes will hurt

Suppose we have 1000 dimensional feature vectors

- But only 10 features are relevant
- Distances will be dominated by the large number of irrelevant features



# Of course, irrelevant attributes will hurt

Suppose we have 1000 dimensional feature vectors

- But only 10 features are relevant
- Distances will be dominated by the large number of irrelevant features

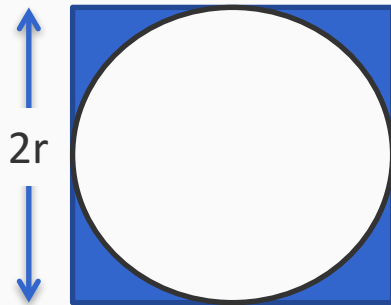
*But even with only relevant attributes, high dimensional spaces behave in odd ways*

# The Curse of Dimensionality

Intuitions that are based on 2 or 3 dimensional spaces do not always carry over to high dimensional spaces

*Example 1:* What fraction of the points in a cube lie outside the sphere inscribed in it?

In two dimensions



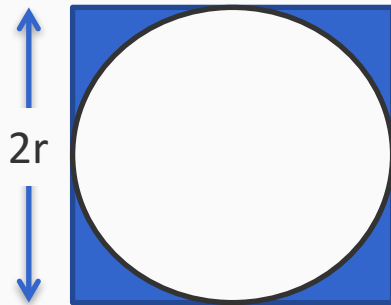
What fraction of the square (i.e the cube) is outside the inscribed circle (i.e the sphere) in two dimensions?

# The Curse of Dimensionality

Intuitions that are based on 2 or 3 dimensional spaces do not always carry over to high dimensional spaces

*Example 1:* What fraction of the points in a cube lie outside the sphere inscribed in it?

In two dimensions



What fraction of the square (i.e the cube) is outside the inscribed circle (i.e the sphere) in two dimensions?

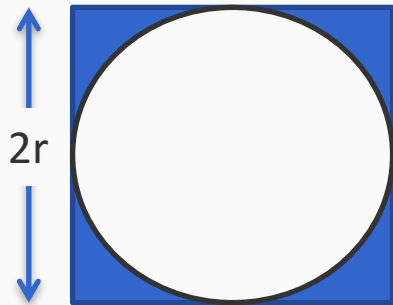
$$1 - \frac{\pi r^2}{4r^2} = 1 - \frac{\pi}{4}$$

# The Curse of Dimensionality

Intuitions that are based on 2 or 3 dimensional spaces do not always carry over to high dimensional spaces

*Example 1:* What fraction of the points in a cube lie outside the sphere inscribed in it?

In two dimensions



What fraction of the square (i.e the cube) is outside the inscribed circle (i.e the sphere) in two dimensions?

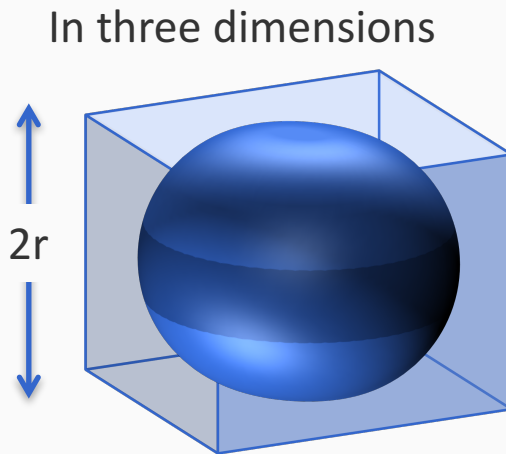
$$1 - \frac{\pi r^2}{4r^2} = 1 - \frac{\pi}{4}$$

But, distances do not behave the same way in high dimensions

# The Curse of Dimensionality

Intuitions that are based on 2 or 3 dimensional spaces do not always carry over to high dimensional spaces

*Example 1:* What fraction of the points in a cube lie outside the sphere inscribed in it?



What fraction of the cube is outside the inscribed sphere in three dimensions?

$$1 - \frac{\frac{4}{3}\pi r^3}{8r^3} = 1 - \frac{\pi}{6}$$

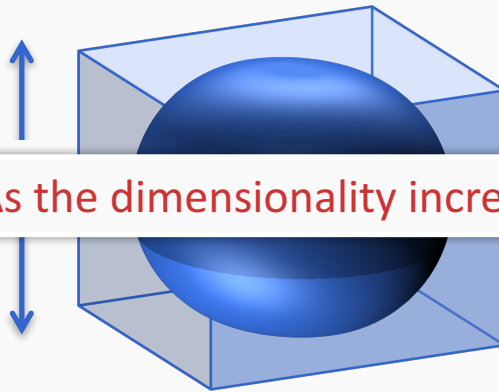
But, distances do not behave the same way in high dimensions

# The Curse of Dimensionality

Intuitions that are based on 2 or 3 dimensional spaces do not always carry over to high dimensional spaces

*Example 1:* What fraction of the points in a cube lie outside the sphere inscribed in it?

In three dimensions



What fraction of the cube is outside the inscribed sphere in three dimensions?

$$\frac{4\pi r^3}{3}$$

As the dimensionality increases, this fraction approaches 1!!

$$\frac{4\pi r^3}{3}$$

In high dimensions, most of the volume of the cube is far away from the center!

# The Curse of Dimensionality

Intuitions that are based on 2 or 3 dimensional spaces do not always carry over to high dimensional spaces

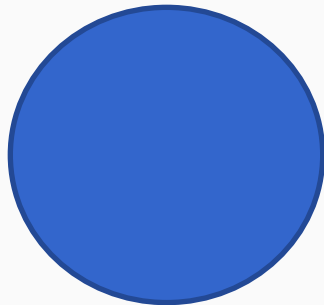
*Example 2:* What fraction of the volume of a unit sphere lies between radius  $1 - \epsilon$  and radius 1?

# The Curse of Dimensionality

Intuitions that are based on 2 or 3 dimensional spaces do not always carry over to high dimensional spaces

*Example 2:* What fraction of the volume of a unit sphere lies between radius  $1 - \epsilon$  and radius 1?

In two dimensions



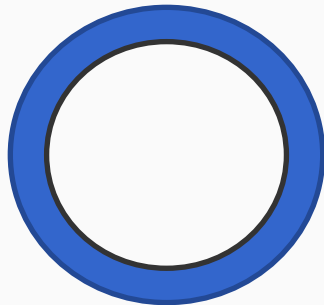


# The Curse of Dimensionality

Intuitions that are based on 2 or 3 dimensional spaces do not always carry over to high dimensional spaces

*Example 2:* What fraction of the volume of a unit sphere lies between radius  $1 - \epsilon$  and radius 1?

In two dimensions



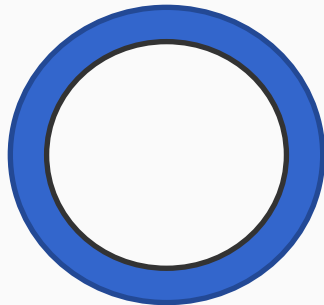
What fraction of the area of the circle is in the blue region?

# The Curse of Dimensionality

Intuitions that are based on 2 or 3 dimensional spaces do not always carry over to high dimensional spaces

*Example 2:* What fraction of the volume of a unit sphere lies between radius  $1 - \epsilon$  and radius 1?

In two dimensions



What fraction of the area of the circle is in the blue region?

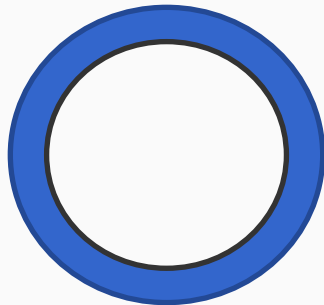
$$\frac{\pi \cdot 1^2 - \pi(1 - \epsilon)^2}{\pi \cdot 1^2} = 1 - (1 - \epsilon)^2$$

# The Curse of Dimensionality

Intuitions that are based on 2 or 3 dimensional spaces do not always carry over to high dimensional spaces

*Example 2:* What fraction of the volume of a unit sphere lies between radius  $1 - \epsilon$  and radius 1?

In two dimensions



What fraction of the area of the circle is in the blue region?

$$\frac{\pi \cdot 1^2 - \pi(1 - \epsilon)^2}{\pi \cdot 1^2} = 1 - (1 - \epsilon)^2$$

But, distances do not behave the same way in high dimensions

# The Curse of Dimensionality

Intuitions that are based on 2 or 3 dimensional spaces do not always carry over to high dimensional spaces

*Example 2:* What fraction of the volume of a unit sphere lies between radius  $1 - \epsilon$  and radius 1?

But, distances do not behave the same way in high dimensions

In  $d$  dimensions, the fraction is  $1 - (1 - \epsilon)^d$

As  $d$  increases, this fraction goes to 1!

In high dimensions, most of the volume of the sphere is far away from the center!

Questions?

# The Curse of Dimensionality

- Most of the points in high dimensional spaces are far away from the origin!
  - In 2 or 3 dimensions, most points are near the center
  - Need more data to “fill up the space”
- Bad news for nearest neighbor classification in high dimensional spaces

Even if most/all features are relevant, in high dimensional spaces, most points are equally far from each other!

“Neighborhood” becomes very large

Presents computational problems too

# Dealing with the curse of dimensionality

- Most “*real-world*” data is not uniformly distributed in the high dimensional space
  - Different ways of capturing the *underlying dimensionality* of the space
  - We will see dimensionality reduction later in the semester
- **Feature selection**, an art
  - Different methods exist
  - Select features, maybe by information gain
  - Try out different feature sets of different sizes and pick a good set based on a validation set
- Prior knowledge or preferences about the hypotheses can also help

Questions?

# Summary: Nearest neighbors classification

- Probably the oldest and simplest learning algorithm
  - Prediction is expensive.
    - Efficient data structures help. **k-D trees**: the most popular, works well in low dimensions
    - Approximate nearest neighbors may be good enough some times. Hashing based algorithms exist
- Requires a distance measure between instances
  - Metric learning: Learn the “right” distance for your problem
- Partitions the space into a Voronoi Diagram
- Beware the curse of dimensionality

Questions?

# Exercises

1. What will happen when you choose  $K$  to the number of training examples?
2. Suppose you want to build a nearest neighbors classifier to predict whether a beverage is a coffee or a tea using two features: the volume of the liquid (in milliliters) and the caffeine content (in grams). You collect the following data:

| Volume<br>(ml) | Caffeine<br>(g) | Label  |
|----------------|-----------------|--------|
| 238            | 0.026           | Tea    |
| 100            | 0.011           | Tea    |
| 120            | 0.040           | Coffee |
| 237            | 0.095           | Coffee |

What is the label for a test point with Volume = 120, Caffeine = 0.013?

Why might this be incorrect?

How would you fix the problem?



# Exercises

1. What will happen when you choose  $K$  to the number of training examples? *The label will always be the most common label in the training data*
2. Suppose you want to build a nearest neighbors classifier to predict whether a beverage is a coffee or a tea using two features: the volume of the liquid (in milliliters) and the caffeine content (in grams). You collect the following data:

| Volume (ml) | Caffeine (g) | Label  |
|-------------|--------------|--------|
| 238         | 0.026        | Tea    |
| 100         | 0.011        | Tea    |
| 120         | 0.040        | Coffee |
| 237         | 0.095        | Coffee |

What is the label for a test point with Volume = 120, Caffeine = 0.013?

*Coffee*

Why might this be incorrect?

*Because Volume will dominate the distance*

How would you fix the problem?

*Rescale the features. Maybe to zero mean, unit variance*